



# Reinforcement Learning

Amr Mohamed Nasr (47)

Michael Raafat (57)

Mohammed Deifallah (59)



## OVERVIEW

For this assignment, you will be making a *Maze Solver*. Your program will generate a maze of size  $N \times N$ . Also, you should generate barriers at random grid locations. Then you will try to learn the path out of the grid using **Policy Iteration** and **Value Iteration**.

## ALGORITHMS USED

As mentioned above, the assignment has two main algorithms to be implemented and applied to reach The objective. Below are more details about them:

### 1) POLICY ITERATION

#### I. ALGORITHM

### Policy Iteration (PI)

1.  $i=0$ ; Initialize  $\pi_0(s)$  randomly for all states  $s$

2. While  $i \neq 0$  or  $|\pi_i - \pi_{i-1}| > 0$  ←

- Policy **evaluation**: Compute value of  $\pi_i$

- $i=i+1$

- Policy **improvement**:

$$Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V^{\pi_i}(s')$$

$$\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$$

Use a L1 norm:  
measures if the policy changed for any state

#### II. DATA STRUCTURE

It only uses 2D matrices to represent the maze cells with corresponding values and actions. Also, it uses a chronologically-ordered list to keep the steps for tracing issues.

## 2) VALUE ITERATION

### I. ALGORITHM

#### Value Iteration (VI)

1. Initialize  $V_0(s)=0$  for all states  $s$
2. Set  $k=1$
3. Loop until [finite horizon, convergence]

- For each state  $s$

$$V_{k+1}(s) = \max_a R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s')$$

- View as Bellman backup on value function

$$V_{k+1} = BV_k$$
$$\pi_{k+1}(s) = \arg \max_a R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s')$$

### II. DATA STRUCTURE

Again, it works with 2D matrices to represent the system.

## CODE ORGANIZATION

The code is mainly divided into three main components:

### 1) ENVIRONMENT

#### 1. CELL

```
private int r, c;  
private Type cell_type;  
private Subtype cell_subtype;  
private List<Action> possible_actions;
```

It has the necessary attributes to represent each state.

Also, note that each cell has a type and subtype. The following figures show the possible categories of them:

```
public enum Type {  
    START, OPEN, BARRIER, END;  
}
```

```
public enum Subtype {  
    MIDDLE, EDGE, CORNER;  
}
```

#### 2. GRID

It has a 2D matrix of **Cells** and the size of the grid world.

Note that this code snippet shows how the probability is being applied to decide the barrier cells.

```

private void initGrid(double r) {
    Random random = new Random();
    for (int i = 0; i < size; i++) {
        for (int j = 0; j < size; j++) {
            if (grid[i][j].getType() == Type.OPEN) {
                double prob = random.nextDouble();
                if (prob < r) {
                    grid[i][j].setType(Type.BARRIER);
                }
            }
        }
    }
}

```

## 2) BRAIN

### 1. POLICY INTERFACE

```

boolean comparePolicy(Policy policy);

void setRandActions();

void setCells(Cell cell[][]);
void setActions(Action grid[][]);
void setValues(Double values[][]);

Cell[][] getGrid();
Double[][] getValues();
Action[][] getActions();

```

## 2. POLICY EVALUATION

```
while (diff > 1e-3 && num_loops < 1e9) {
    diff = 0.0d;
    num_loops++;
    Double[][] new_state_values = new Double[n][n];
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            Cell cell = grid.getGrid()[i][j];
            Action action = actions[i][j];
            double maxVal = -Double.MAX_VALUE;
            if (action == Action.NONE) {
                maxVal = cell.getReward();
            } else {
                Point nextState = action.getTransition(cell);
                maxVal = cell.getReward() + gamma * state_values[nextState.x][nextState.y];
            }
            diff += Math.abs(state_values[i][j] - maxVal);
            new_state_values[i][j] = maxVal;
        }
    }
    state_values = new_state_values;
}
```

## 3. POLICY IMPROVEMENT

The following code snippet shows the application of **Greedy Policy** to improve the current policy:

```

Policy getGreedyPolicy(Double[][] vals, Grid grid, Double gamma) {
    Policy pol = new PolicyImp(grid.getGrid());
    pol.setValues(vals);
    Action actions[][] = new Action[grid.getSize()][grid.getSize()];
    int n = grid.getSize();
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            actions[i][j] = Action.NONE;
            Cell cell = grid.getGrid()[i][j];
            double maxVal = -Double.MAX_VALUE;
            for (int k = 0; k < cell.getPossibleActions().size(); k++) {
                Point nextState = cell.getPossibleActions().get(k).getTransition(cell);
                double val = cell.getReward() + gamma * vals[nextState.x][nextState.y];
                if (val > maxVal) {
                    maxVal = val;
                    actions[i][j] = cell.getPossibleActions().get(k);
                }
            }
        }
    }
    pol.setActions(actions);
    return pol;
}

```

#### 4. VALUE INTERFACE

```

public interface ValueIteration {

    void value_iteration(Grid grid);
    List<Double[][]> getValueIterations();
    Policy getOptimalPolicy();
}

```

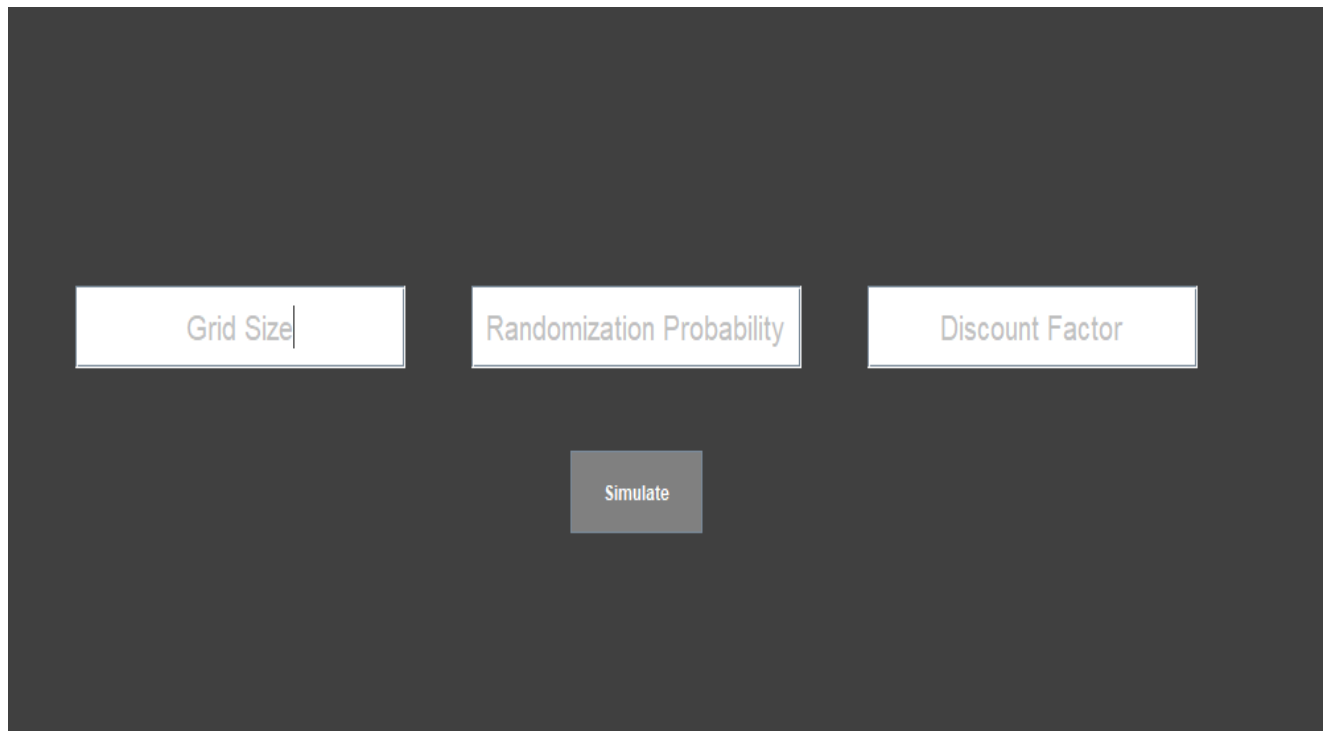
#### 3) GUI

It's a separate package for our visualization.

#### EXTRA WORK

- 1) We provide a GUI application, implemented in **Java Swing**, that enables the user to try our solver for different purposes:
  - 1- Measure the time performance.
  - 2- Try the various effects of parameters.

Here's a screenshot of the application intro:



- 2) We provide a separate implementation of DFS function that guarantees path existence between the start and end cells.

```
public static boolean pathExists(Cell grid[][]) {
    List<Point> stack = new ArrayList<Point>();
    stack.add(grid[0][0].getPosition());
    Set<Point> visited = new HashSet<Point>();
    while(!stack.isEmpty()) {
        Point point = stack.remove(stack.size() - 1);
        visited.add(grid[point.x][point.y].getPosition());
        if (grid[point.x][point.y].getType() == Type.END) {
            return true;
        }
        // searching for valid neighbors cells
        List<Action> actions = grid[point.x][point.y].getPossibleActions();
        for (int i = 0; i < actions.size(); i++) {
            Action action = actions.get(i);
            if (!visited.contains(action.getTransition(grid[point.x][point.y]))) {
                stack.add(action.getTransition(grid[point.x][point.y]));
            }
        }
    }
    return false;
}
```





## SAMPLE RUNS

The following sample run has a maze of size 6×6 with probability of .2 and discount factor of .9

S	O	O	B	B	O
O	B	O	O	O	O
O	O	O	O	O	O
O	O	B	O	O	O
B	O	O	O	O	O
O	B	O	O	O	T

Maze

→ → ↓ - - ↓	-6.5132155990000005	-6.12579511	-5.6953279000000006	-	-	-4.0951
↓ - → → → ↓	-6.12579511	-	-5.217031	-4.68559	-4.0951	-3.439
→ → → → → ↓	-5.6953279000000006	-5.217031	-4.68559	-4.0951	-3.439	-2.71
→ ↓ - → → ↓	-5.217031	-4.68559	-	-3.439	-2.71	-1.9
- → → → → ↓	-	-4.0951	-3.439	-2.71	-1.9	-1.0
→ - → → → -	-1.9327352842E9	-	-2.71	-1.9	-1.0	0.0

Value Iteration  
Iteration 11 of 11  
Time taken : 2 ms

trace

Value Iteration

→ → ↓ - - ↓	-6.5132155990000005	-6.12579511	-5.6953279000000006	-	-	-4.0951
↓ - → → → ↓	-6.12579511	-	-5.217031	-4.68559	-4.0951	-3.439
→ → → → → ↓	-5.6953279000000006	-5.217031	-4.68559	-4.0951	-3.439	-2.71
→ ↓ - → → ↓	-5.217031	-4.68559	-	-3.439	-2.71	-1.9
- → → → → ↓	-	-4.0951	-3.439	-2.71	-1.9	-1.0
→ - → → → -	-1.9327352842E9	-	-2.71	-1.9	-1.0	0.0

Policy Iteration  
Iteration 5 of 5  
Time taken : 10 ms

trace

Policy Iteration

## ASSUMPTIONS & EXPLANATIONS

- The following table shows the corresponding rewards for each type of cell:

Cell Type	Reward
Start	-1
Middle	-1
Barrier	-2147483648 (minimum integer in Java)
End	0

- We only permit actions that have effect. In other words, each corner cell has only 2 actions, each edge cell has 3 actions and any other middle cell has the 4 default actions.
- The program has 3 inputs:
  1. N: The grid world size.
  2. P: the probability of barriers.  $0 \leq P \leq 1$
  3.  $\gamma$ : The discount factor  $0 \leq \gamma < 1$

Note that the discount factor can't be 1 to prevent possible infinite loops in **Policy Iteration**.

- The algorithms terminate if any of the following conditions holds:
  1. The difference between two consecutive iterations is less than .001
  2. The number of iterations has reached  $10^9$
- The recommended configuration of the maze to really measure the performance is: {N: 8, P: 0.2,  $\gamma$ : .9}